

Datasets

This list of datasets below has been compiled for the convenience of the students for reference purposes only. The list is provided as it is, without warranty of any kind. It is students' responsibility to wisely select and carefully check the datasets according to the guidelines and discussion outlined in the lectures prior to use. Note that the list below is not exhaustive, and it is the responsibility of the students to conduct thorough research to identify (additional) suitable datasets for the Group Mini Project.

Please note that some datasets may only be obtainable upon request from the dataset authors. In these cases, it is recommended to contact the authors as soon as possible to avoid any delays this process might create.

HateMM Dataset

HateMM: A Multi-Modal Dataset for Hate Video Classification (2023)

Paper: <https://arxiv.org/abs/2305.03915v1>

Repository: <https://zenodo.org/records/7799469>

Creative Commons Attribution 4.0 International

OLID: Offensive Language Identification Dataset

Predicting the Type and Target of Offensive Posts in Social Media (2020)

Paper: <https://arxiv.org/pdf/1902.09666v2>

Repository:

CC0, with the underlying comment text being governed by Wikipedia's CC-SA-3.0.

Hate Speech and Offensive Language Dataset

Automated Hate Speech Detection and the Problem of Offensive Language (2017)

Paper: <https://ojs.aaai.org/index.php/ICWSM/article/view/14955/14805>

Repository: <https://github.com/t-davidson/hate-speech-and-offensive-language>

MIT License

ThreatGram 101: Extreme Telegram Replies Data with Threat Levels Dataset

Exploring Multi-Level Threats in Telegram Data with AI-Human Annotation: A Preliminary Study (2023)

Paper: <https://ieeexplore.ieee.org/document/10459792>

Repository: <https://data.mendeley.com/datasets/tm9s68vgxd/1>

CC BY 4.0

Implicit Hate Dataset

Latent Hatred: A Benchmark for Understanding Implicit Hate Speech (2021)

Paper: <https://aclanthology.org/2021.emnlp-main.29.pdf>

Repository: <https://github.com/SALT-NLP/implicit-hate>

MIT License

Civil Comments: Jigsaw Unintended Bias in Toxicity Classification Dataset

Nuanced Metrics for Measuring Unintended Bias with Real Data for Text Classification (2019)

Paper: <https://dl.acm.org/doi/10.1145/3308560.3317593>

Repository: https://huggingface.co/datasets/google/civil_comments

Creative Commons Zero v1.0 Universal

Hatemoji Dataset

Hatemoji: A Test Suite and Adversarially-Generated Dataset for Benchmarking and Detecting Emoji-based Hate (2022)

Paper: <https://arxiv.org/abs/2108.05921>

Repository: <https://github.com/HannahKirk/Hatemoji>

CC-BY-4.0 license

More datasets are available (but not limited to) at: <https://github.com/leondz/hatespeechdata/tree/master>

HatEval Dataset

SemEval-2019 Task 5: Multilingual Detection of Hate Speech Against Immigrants and Women in Twitter (2019)

Paper: <https://aclanthology.org/S19-2007.pdf>

Repository: <https://github.com/msang/hateval>

Creative Commons CC-BY-NC-4.0

MetaHate Dataset

MetaHate: A Dataset for Unifying Efforts on Hate Speech Detection (2024)

Paper: <https://ojs.aaai.org/index.php/ICWSM/article/view/31445/33605>

Repository: <https://github.com/palomapiot/metahate?tab=readme-ov-file>

Creative Commons Attribution Non Commercial Share Alike 4.0

HateSpeech Dataset

Hate Speech Dataset from a White Supremacy Forum (2018)

Paper: <https://aclanthology.org/W18-51.pdf>

Repository: <https://github.com/Vicomtech/hate-speech-dataset?tab=License-1-ov-file>

Creative Commons Attribution-ShareAlike 3.0 Spain License

ETHOS: multi-label haTe speech detectiOn dataSet

ETHOS: an Online Hate Speech Detection Dataset (2021)

Paper: <https://link.springer.com/epdf/10.1007/s40747-021-00608-2>

Repository: <https://github.com/intelligence-csd-auth-gr/Ethos-Hate-Speech-Dataset>

GNU GPLv3

MLMA Hate Speech Dataset

Multilingual and Multi-Aspect Hate Speech Analysis (2019)

Paper: <https://aclanthology.org/D19-1474.pdf>

Repository: https://github.com/HKUST-KnowComp/MLMA_hate_speech

MIT license

CONAN: COUNTER NARRATIVES THROUGH NICHESOURCING DATASET

CONAN - COUNTER NARRATIVES THROUGH NICHESOURCING: A MULTILINGUAL DATASET OF RESPONSES TO FIGHT ONLINE HATE SPEECH (2019)

Paper: <https://aclanthology.org/P19-1271.pdf>

Repository: <https://github.com/marcoguerini/CONAN>

“These resources can be used for research purposes and cannot be redistributed. Please cite the corresponding publication if you use any dataset.”

Knowledge-grounded hate countering Dataset

Towards Knowledge-Grounded Counter Narrative Generation for Hate Speech (2021)

Paper: <https://aclanthology.org/2021.findings-acl.79.pdf>

Repository: <https://github.com/marcoguerini/CONAN>

“These resources can be used for research purposes and cannot be redistributed. Please cite the corresponding publication if you use any dataset.”

Multitarget CONAN Dataset

Human-in-the-Loop for Data Collection: a Multi-Target Counter Narrative Dataset to Fight Online Hate Speech (2021)

Paper: <https://aclanthology.org/2021.acl-long.250.pdf>

Repository: <https://github.com/marcoquerini/CONAN>

“These resources can be used for research purposes and cannot be redistributed. Please cite the corresponding publication if you use any dataset.”

DIALOCONAN Dataset

Human-Machine Collaboration Approaches to Build a Dialogue Dataset for Hate Speech Countering (2022)

Paper: <https://aclanthology.org/2022.emnlp-main.549.pdf>

Repository: <https://github.com/marcoquerini/CONAN>

“These resources can be used for research purposes and cannot be redistributed. Please cite the corresponding publication if you use any dataset.”

Hate Speech and Offensive Language Dataset

Automated Hate Speech Detection and the Problem of Offensive Language (2017)

Paper: <https://ojs.aaai.org/index.php/ICWSM/article/view/14955/14805>

Repository: <https://github.com/t-davidson/hate-speech-and-offensive-language>

Note: the repository is no longer actively maintained.

MIT License

TOXIGEN Dataset

ToxiGen: A Large-Scale Machine-Generated Dataset for Adversarial and Implicit Hate Speech Detection (2022)

Paper: <https://arxiv.org/abs/2203.09509>

Repository: <https://github.com/microsoft/ToxiGen>

Multimodal Sarcasm Detection Dataset: MUsTARD Dataset

Towards Multimodal Sarcasm Detection (An Obviously Perfect Paper) (2019)

Note: The dataset is compiled from popular TV shows including Friends, The Golden Girls, The Big Bang Theory, and Sarcasmaholics Anonymous.

Paper: <https://aclanthology.org/P19-1455/>

Repository: <https://github.com/soujanyaoria/MUsTARD>

MUsTARD++ Dataset

A Multimodal Corpus for Emotion Recognition in Sarcasm (2022)

Note: The dataset is compiled from popular TV shows including Friends, The Golden Girls, The Big Bang Theory, and Sarcasmaholics Anonymous.

Paper: <https://arxiv.org/abs/2206.02119>

Repository: https://github.com/cfiltnlp/MUsTARD_Plus_Plus

Online Misogyny Dataset

An expert annotated dataset for the detection of online misogyny (2021)

Paper: <https://aclanthology.org/2021.eacl-main.114/>

Repository: <https://github.com/ellamquest/online-misogyny-eacl2021>

CAD: Contextual Abuse Dataset

Introducing CAD: the Contextual Abuse Dataset (2021)

Paper: <https://aclanthology.org/2021.naacl-main.182/>

Repository: https://github.com/dongpng/cad_naacl2021, <https://zenodo.org/records/4881008>

A Multilingual Dataset of Racial Stereotypes Dataset

A Multilingual Dataset of Racial Stereotypes in Social Media Conversational Threads (2023)

Paper: <https://aclanthology.org/2023.findings-eacl.51/>

Note: The annotated dataset will be available for research purposes upon request, together with the complete set of annotation guidelines.

News-Headlines-Dataset-For-Sarcasm-Detection

Sarcasm detection using news headlines Dataset (2023)

Paper: <https://www.sciencedirect.com/science/article/pii/S2666651023000013#fn5>

Repository: <https://github.com/rishabhmisra/News-Headlines-Dataset-For-Sarcasm-Detection>
<https://www.kaggle.com/datasets/rmisra/news-headlines-dataset-for-sarcasm-detection/data>

iSarcasm Dataset

iSarcasm: A Dataset of Intended Sarcasm (2019)

Paper: <https://arxiv.org/abs/1911.03123v2>

Repository: <https://github.com/dmbavkar/iSarcasm>

Note from the dataset authors: While we only make the tweet IDs public, we maintain and are happy to provide the following for research purposes, under an agreement that protects the privacy of our users: tweet texts; for each sarcastic tweet, an explanation given by its authors as to why the tweet is sarcastic; for each sarcastic tweet, a rephrase given by its author that conveys the same message non-sarcastically.

Large-Scale Hate Speech Detection Dataset

Large-Scale Hate Speech Detection with Cross-Domain Transfer (2022)

Paper: <https://aclanthology.org/2022.lrec-1.238/>

Repository: <https://github.com/avaapm/hatespeech>

HateXplain Dataset

HateXplain: A Benchmark Dataset for Explainable Hate Speech Detection (2021)

Paper: <https://ojs.aaai.org/index.php/AAAI/article/view/17745>

Repository: <https://github.com/hate-alert/HateXplain>

ConvAbuse Dataset

ConvAbuse: Data, Analysis, and Benchmarks for Nuanced Abuse Detection in Conversational AI (2021)

Paper: <https://aclanthology.org/2021.emnlp-main.587/>

Repository: <https://github.com/amandacurry/convabuse>

SWAD: Swear Words Abusiveness Dataset

Do You Really Want to Hurt Me? Predicting Abusive Swearing in Social Media (2020)

Paper: <https://aclanthology.org/2020.lrec-1.765/>

Repository: <https://github.com/dadangewp/SWAD-Repository>

World of Warcraft Cyberbullying Dataset

Detecting cyberbullying in online communities (2016)

Paper: <https://core.ac.uk/download/pdf/301369744.pdf>

Repository: <http://ub-web.de/research/>

League of Legends Cyberbullying Dataset

Detecting cyberbullying in online communities (2016)

Paper: <https://core.ac.uk/download/pdf/301369744.pdf>

Repository: <http://ub-web.de/research/>

MultiOFF: Multimodal Meme Dataset

Multimodal Meme Dataset (MultiOFF) for Identifying Offensive Content in Image and Text (2020)

Paper: <https://aclanthology.org/2020.trac-1.6.pdf>

Repository: <https://drive.google.com/drive/folders/1hKLOtpVmF45loBmJPwojgg6XraLtHmV6>

ALONE: AdoLescents ON twittEr Dataset

ALONE: A Dataset for Toxic Behavior among Adolescents on Twitter (2020)

Paper: <https://arxiv.org/pdf/2008.06465.pdf>

Repository: <https://zenodo.org/records/3608352>

Note from the dataset authors: we make this dataset available upon request to the authors, and researchers will be required to sign an agreement to use it only for research purposes and without public dissemination.

MMHS150K Dataset

Exploring Hate Speech Detection in Multimodal Publications (2019)

Paper: <https://arxiv.org/pdf/1910.03814.pdf>

Repository: <https://gombbru.github.io/2019/10/09/MMHS/>

HarMeme Dataset

Detecting Harmful Memes and Their Targets (2021)

Paper: <https://arxiv.org/abs/2110.00413>

Repository: <https://github.com/di-dimitrov/harmeme>

OASST1: OpenAssistant Conversations Dataset

OpenAssistant Conversations – Democratizing Large Language Model Alignment (2023)

Paper: https://proceedings.neurips.cc/paper_files/paper/2023/file/949f0f8f32267d297c2d4e3ee10a2e7e-Paper-Datasets_and_Benchmarks.pdf

Repository: <https://github.com/LAION-AI/Open-Assistant>

Apache-2.0 license

FairPrism Dataset

FairPrism: Evaluating Fairness-Related Harms in Text Generation (2023)

Repository: <https://github.com/microsoft/FairPrism>

Paper: <https://aclanthology.org/2023.acl-long.343.pdf>

MIT license

Note on dataset access: reach out to the fairprism@microsoft.com to access the dataset.

HarmfulQA Dataset

Red-Teaming Large Language Models using Chain of Utterances for Safety-Alignment (2023)

Paper: <https://arxiv.org/abs/2308.09662>

Repository: <https://github.com/declare-lab/red-instruct>

Apache-2.0 license

MultiHateClip Dataset

MultiHateClip: A Multilingual Benchmark Dataset for Hateful Video Detection on YouTube and Bilibili (2024)

Paper: <https://dl.acm.org/doi/pdf/10.1145/3664647.3681521>

Repository: <https://github.com/social-ai-studio/multihateclip?tab=readme-ov-file>

MAMI: Multimedia Automatic Misogyny Identification Dataset

SemEval-2022 Task 5: Multimedia Automatic Misogyny Identification (2022)

Paper: <https://aclanthology.org/2022.semeval-1.74/>

Repository: <https://github.com/MIND-Lab/SemEval2022-Task-5-Multimedia-Automatic-Misogyny-Identification-MAMI->

“The data may be distributed upon request and for academic purposes only. To request the datasets, please fill out the following form: <https://forms.gle/AGWMMiGicBHiQx4q98>

After submitting the required info, participants will have a link to a folder containing the datasets in a zip format (trial, training and development) and the password to uncompress the files.”

HateComments Dataset

Hateful Comment Detection and Hate Target-Type Prediction for Video Comments (2023)

Paper: <https://dl.acm.org/doi/pdf/10.1145/3583780.3615260>

Repository: <https://drive.google.com/file/d/1EUbWUokv1CYkWKlwByUC6yIuBGUw2MN/>

The “Call me sexist, but” sexism dataset

Paper: <https://ojs.aaai.org/index.php/ICWSM/article/view/18085/17888>

Repository: https://search.gesis.org/research_data/SDN-10.7802-2251?doi=10.7802

Website flickers

Hateful Meme Dataset

The Hateful Memes Challenge: Detecting Hate Speech in Multimodal Memes (2020)

Paper: <https://arxiv.org/abs/2005.04790>

Repository: https://github.com/facebookresearch/mmf/tree/main/projects/hateful_memes
<https://ai.meta.com/tools/hatefulmemes/>

To acquire the data, you will need to register at DrivenData's Hateful Memes Competition and download data from the challenge's [download page](#). MMF won't be able to automatically download the data since you manually need to agree to the licensing terms.

BSD License